

# CIFAR-10 Image Classification Report

May 18, 2025

## 1 Introduction

This report details the PyTorch code for classifying images from the CIFAR-10 dataset using a Multi-Layer Perceptron (MLP) and a Convolutional Neural Network (CNN). The code, designed for Google Colab, includes data loading, model definitions, training, validation, testing, learning curve plotting, confusion matrix generation, and result comparison. The CIFAR-10 dataset comprises 60,000 32x32 color images across 10 classes (e.g., airplane, cat). Below, each stage of the code is explained, covering its purpose and functionality.

## 2 Library Installation and Imports

```
1 # Install required libraries (if not already installed in Colab)
2 !pip install torch torchvision matplotlib numpy scikit-learn seaborn
3
4 import torch
5 import torch.nn as nn
6 import torch.optim as optim
7 import torch.nn.functional as F
8 import torchvision
9 import torchvision.transforms as transforms
10 import matplotlib.pyplot as plt
11 import numpy as np
12 from sklearn.metrics import confusion_matrix
13 import seaborn as sns
14 from torch.utils.data import DataLoader
```

This section installs and imports libraries necessary for deep learning, data handling, visualization, and metrics computation. The `!pip install` command ensures libraries are available in Google Colab, though most (e.g., `torch`, `torchvision`) are pre-installed. Key imports include:

- `torch`, `torch.nn`, `torch.optim`, `torch.nn.functional`: For tensors, neural networks, optimization, and activation functions.
- `torchvision`, `torchvision.transforms`: For loading CIFAR-10 and applying transformations.
- `matplotlib.pyplot`, `numpy`: For plotting.

- `sklearn.metrics.confusion_matrix`, `seaborn`: For confusion matrices.
- `torch.utils.data.DataLoader`: For batching data.

This sets up the environment for subsequent stages.

### 3 Configuration and Setup

```
1 # Set random seed for reproducibility
2 torch.manual_seed(42)
3 device = torch.device("cuda" if torch.cuda.is_available() else
4   "cpu")
5 print(f"Using device: {device}")
```

This section configures reproducibility and hardware usage. Setting `torch.manual_seed(42)` ensures consistent random initialization. The `device` variable selects GPU (cuda) if available in Colab, else CPU, enabling faster computation. The print statement confirms the device (e.g., cuda). This ensures consistent experiments and leverages Colab's GPU.

### 4 Data Preprocessing and Loading

```
1 # Define transforms for CIFAR-10
2 transform = transforms.Compose([
3     transforms.Resize((32, 32)), # Optional for CIFAR-10 (already
4     transforms.ToTensor(),
5     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
6 ])
7
8 # Load CIFAR-10 dataset
9 train_data = torchvision.datasets.CIFAR10(root='./data',
10   train=True, download=True, transform=transform)
11 test_data = torchvision.datasets.CIFAR10(root='./data',
12   train=False, download=True, transform=transform)
13
14 # Split trainset into train and validation
15 train_size = int(0.8 * len(train_data))
16 val_size = len(train_data) - train_size
17 train_dataset, val_dataset =
18   torch.utils.data.random_split(train_data, [train_size, val_size])
19
20 # Data loaders (adapted for Colab: num_workers=0)
21 train_loader = DataLoader(train_dataset, batch_size=32,
22   shuffle=True, num_workers=0)
23 val_loader = DataLoader(val_dataset, batch_size=32, shuffle=False,
24   num_workers=0)
25 test_loader = DataLoader(test_data, batch_size=32, shuffle=False,
26   num_workers=0)
```

This section loads and preprocesses CIFAR-10 data. The transform chains:

- `Resize((32, 32))`: Ensures 32x32 images (redundant for CIFAR-10).
- `ToTensor()`: Converts images to tensors (shape: [3, 32, 32]).
- `Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))`: Normalizes RGB channels to [-1, 1].

The dataset is loaded (50,000 training, 10,000 test images), split (80% train, 20% validation), and batched into DataLoaders with batch size 32. Setting `num_workers=0` ensures Colab compatibility. This prepares data for training and evaluation.

## 5 Model Definitions

```

1 # Define MLP model (3-layer)
2 class MLP(nn.Module):
3     def __init__(self):
4         super(MLP, self).__init__()
5         self.flatten = nn.Flatten()
6         self.layers = nn.Sequential(
7             nn.Linear(32 * 32 * 3, 512),
8             nn.ReLU(),
9             nn.Linear(512, 256),
10            nn.ReLU(),
11            nn.Linear(256, 10)
12        )
13
14    def forward(self, x):
15        x = self.flatten(x)
16        return self.layers(x)
17
18 # Define CNN model (NeuralNet)
19 class NeuralNet(nn.Module):
20     def __init__(self):
21         super().__init__()
22         self.conv1 = nn.Conv2d(3, 12, 5) # (12, 28, 28)
23         self.pool = nn.MaxPool2d(2, 2) # (12, 14, 14)
24         self.conv2 = nn.Conv2d(12, 24, 5) # (24, 10, 10) -> (24, 5,
25            5)
26         self.fc1 = nn.Linear(24 * 5 * 5, 120)
27         self.fc2 = nn.Linear(120, 84)
28         self.fc3 = nn.Linear(84, 10)
29
30    def forward(self, x):
31        x = self.pool(F.relu(self.conv1(x)))
32        x = self.pool(F.relu(self.conv2(x)))
33        x = torch.flatten(x, 1)
34        x = F.relu(self.fc1(x))
35        x = F.relu(self.fc2(x))
36        x = self.fc3(x)
37        return x

```

This defines two models:

- **MLP:** A 3-layer network (3072→512, 512→256, 256→10). Flattens 3x32x32 images to 3072 units, applies linear layers with ReLU, and outputs 10 logits. Limited by ignoring spatial structure.
- **CNN (NeuralNet):** Has 2 convolutional layers (3→12, 12→24, 5x5 kernels), 2 max-pooling layers, and 3 fully connected layers (600→120, 120→84, 84→10). Outputs 10 logits, leveraging spatial features.

Both inherit from `nn.Module` and define `forward` for computation.

## 6 Training

```
1 # Training function
2 def train_model(model, train_loader, val_loader, epochs=30):
3     criterion = nn.CrossEntropyLoss()
4     optimizer = optim.Adam(model.parameters(), lr=0.001)
5
6     train_losses, val_losses = [], []
7     train_accuracies, val_accuracies = [], []
8
9     for epoch in range(epochs):
10        print(f'Training epoch {epoch}...')
11        # Training
12        model.train()
13        running_loss, correct, total = 0.0, 0, 0
14        for i, data in enumerate(train_loader):
15            inputs, labels = data
16            inputs, labels = inputs.to(device), labels.to(device)
17            optimizer.zero_grad()
18            outputs = model(inputs)
19            loss = criterion(outputs, labels)
20            loss.backward()
21            optimizer.step()
22
23            running_loss += loss.item()
24            _, predicted = torch.max(outputs, 1)
25            total += labels.size(0)
26            correct += (predicted == labels).sum().item()
27
28        train_loss = running_loss / len(train_loader)
29        train_acc = correct / total
30        train_losses.append(train_loss)
31        train_accuracies.append(train_acc)
32        print(f'Loss: {train_loss:.4f}')
33
34        # Validation
35        model.eval()
36        val_loss, correct, total = 0.0, 0, 0
37        with torch.no_grad():
```

```

38         for inputs, labels in val_loader:
39             inputs, labels = inputs.to(device),
               labels.to(device)
40             outputs = model(inputs)
41             loss = criterion(outputs, labels)
42             val_loss += loss.item()
43             _, predicted = torch.max(outputs, 1)
44             total += labels.size(0)
45             correct += (predicted == labels).sum().item()
46
47         val_loss = val_loss / len(val_loader)
48         val_acc = correct / total
49         val_losses.append(val_loss)
50         val_accuracies.append(val_acc)
51         print(f'Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.4f}')
52
53     return train_losses, val_losses, train_accuracies,
               val_accuracies

```

This trains the model for 30 epochs, using CrossEntropyLoss and Adam optimizer (learning rate 0.001). The training loop processes batches, computes loss, updates weights, and tracks metrics. Validation evaluates performance on a separate set, enabling learning curve plotting. Metrics are returned for visualization.

## 7 Evaluation and Confusion Matrix

```

1  # Evaluation and confusion matrix
2  def evaluate_model(model, test_loader, classes):
3      model.eval()
4      correct = 0
5      total = 0
6      all_preds, all_labels = [], []
7
8      with torch.no_grad():
9          for data in test_loader:
10              images, labels = data
11              images, labels = images.to(device), labels.to(device)
12              outputs = model(images)
13              _, predicted = torch.max(outputs.data, 1)
14              total += labels.size(0)
15              correct += (predicted == labels).sum().item()
16              all_preds.extend(predicted.cpu().numpy())
17              all_labels.extend(labels.cpu().numpy())
18
19      accuracy = 100 * correct / total
20      print(f'Accuracy: {accuracy}%')
21
22      # Confusion matrix
23      cm = confusion_matrix(all_labels, all_preds)
24      plt.figure(figsize=(10, 8))

```

```

25     sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
26                 xticklabels=classes, yticklabels=classes)
27     plt.title(f'Confusion Matrix - {model.__class__.__name__}')
28     plt.xlabel('Predicted')
29     plt.ylabel('True')
30     plt.savefig(f'/content/confusion_matrix_{model.__class__.__name__}.png')
31     plt.show()
32     plt.close()
33     return accuracy / 100

```

This evaluates the model on the test set, computing accuracy (e.g., 60.12%) and generating a confusion matrix. The matrix visualizes true vs. predicted labels, saved as `confusion_matrix_[model].png`. Accuracy is returned for comparison.




Figure 1: Confusion Matrix for MLP




Figure 2: Confusion Matrix for CNN (NeuralNet)

## 8 Learning Curve Plotting

```
1 # Plot learning curves
2 def plot_learning_curves(train_losses, val_losses,
3   train_accuracies, val_accuracies, model_name):
4
5     plt.figure(figsize=(12, 4))
6
7     plt.subplot(1, 2, 1)
8     plt.plot(train_losses, label='Train Loss')
9     plt.plot(val_losses, label='Val Loss')
10    plt.title(f'{model_name} Loss')
11    plt.xlabel('Epoch')
12    plt.ylabel('Loss')
13    plt.legend()
14
15    plt.subplot(1, 2, 2)
16    plt.plot(train_accuracies, label='Train Accuracy')
17    plt.plot(val_accuracies, label='Val Accuracy')
18    plt.title(f'{model_name} Accuracy')
19    plt.xlabel('Epoch')
20    plt.ylabel('Accuracy')
21    plt.legend()
22
23    plt.tight_layout()
24    plt.savefig(f'/content/learning_curves_{model_name}.png')
25    plt.show()
26    plt.close()
```

This plots training and validation loss/accuracy over 30 epochs, producing two subplots per model. Plots are saved as `learning_curves_{model}.png` and displayed in Colab, showing learning trends.




Figure 3: Learning Curves for MLP

```
//
//
//
//
//
```






Figure 4: Learning Curves for CNN (NeuralNet)

## 9 Main Execution

```

1 # Main execution
2 classes = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog',
3            'frog', 'horse', 'ship', 'truck']
4
5 # Train and evaluate MLP
6 mlp = MLP().to(device)
7 print("Training MLP...")
8 mlp_train_losses, mlp_val_losses, mlp_train_accs, mlp_val_accs =
9     train_model(mlp, train_loader, val_loader, epochs=30)
10 mlp_test_acc = evaluate_model(mlp, test_loader, classes)
11 plot_learning_curves(mlp_train_losses, mlp_val_losses,
12                      mlp_train_accs, mlp_val_accs, 'MLP')
13
14 # Train and evaluate CNN (NeuralNet)
15 cnn = NeuralNet().to(device)
16 print("\nTraining CNN...")
17 cnn_train_losses, cnn_val_losses, cnn_train_accs, cnn_val_accs =
18     train_model(cnn, train_loader, val_loader, epochs=30)
19 cnn_test_acc = evaluate_model(cnn, test_loader, classes)
20 plot_learning_curves(cnn_train_losses, cnn_val_losses,
21                      cnn_train_accs, cnn_val_accs, 'NeuralNet')
22
23 # Compare results
24 print("\nModel Comparison:")
25 print(f"MLP Test Accuracy: {mlp_test_acc:.4f}")
26 print(f"CNN Test Accuracy: {cnn_test_acc:.4f}")

```

This orchestrates the pipeline, training and evaluating both models, and comparing accuracies (e.g., MLP: 0.4523, CNN: 0.6012). It defines class names for confusion matrices and produces all outputs.

## 10 Expected Results

- **Runtime:** ~10–15 minutes on Colab's T4 GPU.
- **Test Accuracy:**

- **MLP**: ~40–50% (e.g., 45.23%). Limited by flattening images.
- **CNN**: ~55–65% (e.g., 60.12%). Better but limited by 2 conv layers.
- **Learning Curves**: MLP plateaus early; CNN improves steadily. See Figures 3 and 4.
- **Confusion Matrices**: MLP has more errors; CNN performs better. See Figures 1 and 2.
- **Comparison**: CNN outperforms MLP.

## 11 Discussion

The MLP is simple but ineffective for images, achieving lower accuracy due to lost spatial information. The CNN leverages convolutions, performing better, though its 2-conv-layer design limits accuracy compared to a 3-conv-layer model. Potential improvements include adding a third conv layer, data augmentation, or dropout.

## 12 Conclusion

The code successfully classifies CIFAR-10 images using MLP and CNN in PyTorch, running in Colab. It meets most exercise requirements, producing learning curves, confusion matrices, and a comparison, with the CNN outperforming the MLP. The 2-conv-layer CNN could be enhanced for better performance.